


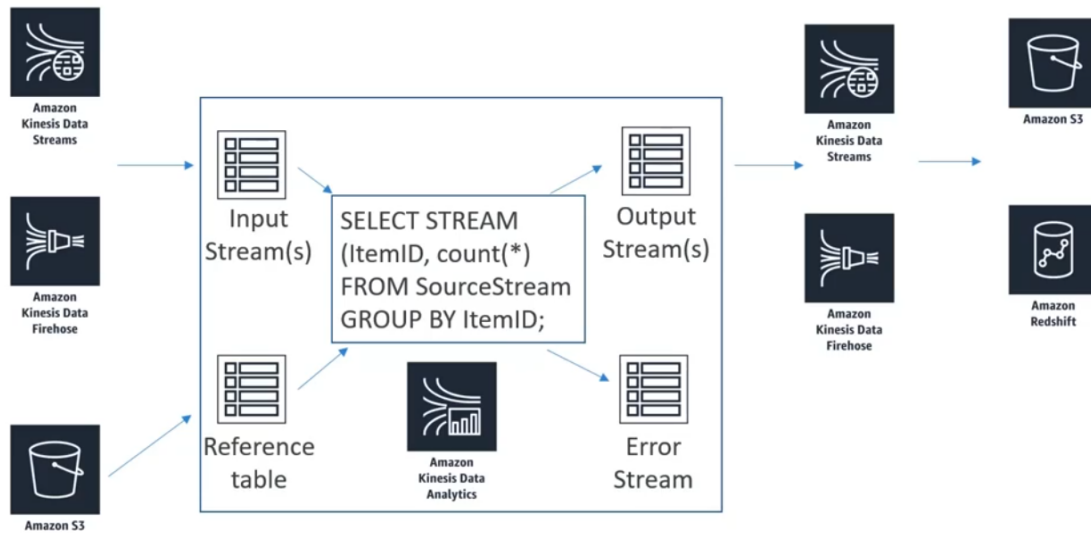
Kinesis Analytics

Querying streams of data

Conceptually...



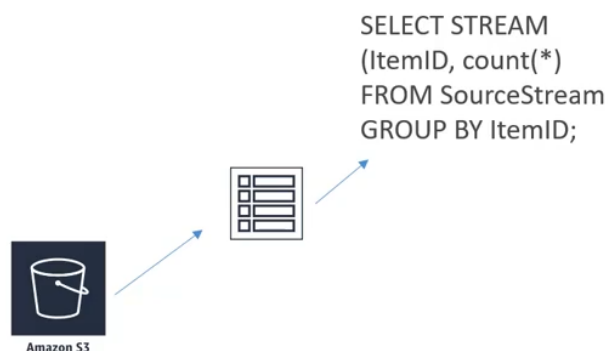
In more depth...



about pic we saw the data is going via amazon kinesis data stream , kinesis data firehouse , amazon s3 bucket is reference table if it got any error stream means broken are some missing values . after processing data it will go to the next process going to kinesis data stream to s3 bucket

Reference tables are cool

- Inexpensive way to “join” data for quick lookups
 - i.e., look up the city associated with a zip code
 - Mapping is stored in S3 which is very inexpensive
 - ↳ Just use a “JOIN” command to use the data in your queries

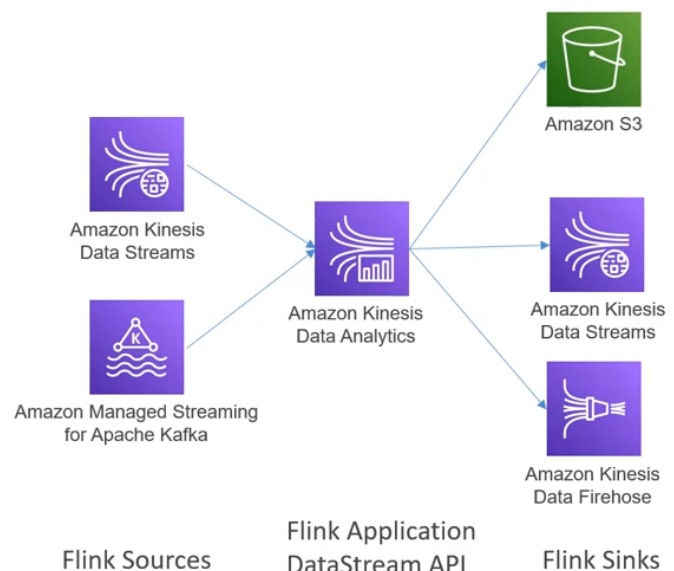


Kinesis Data Analytics + Lambda

- AWS Lambda can be a destination as well
- Allows lots of flexibility for post-processing
 - Aggregating rows
 - Translating to different formats
 - Transforming and enriching data
 - Encryption
- Opens up access to other services & destinations
 - S3, DynamoDB, Aurora, Redshift, SNS, SQS, CloudWatch

Kinesis Data Analytics for Apache Flink

- Formerly Kinesis Data Analytics for Java
 - Kinesis Data Analytics always used Flink under the hood
 - But now supports Scala as well as Java
- Flink is a framework for processing data streams
- Kinesis Data Analytics integrates Flink with AWS
 - Instead of using SQL, you can develop your own Flink application from scratch and load it into KDA via S3
- Serverless



Common use-cases

- Streaming ETL
- Continuous metric generation
- Responsive analytics



Amazon Kinesis Data Analytics

streaming ETL : for example you create an app with continuous IoT sensors in a Kinesis data stream, organise the data, normalise the data, detect if the data was repeated or not, and send the data to an S3 bucket.

continuous metric generation ; for example you create a game leader board to continuously update the first and some use cases.

read the data from Kinesis data stream and analyse the data in Kinesis data analytics, then data will be written to DynamoDB.

check the website unique users for every five minutes and write the data in Redshift for later analysis.

responsive analytics : check the API success rate for payment request and send data to CloudWatch.

Kinesis Analytics

- Pay only for resources consumed (but it's not cheap)
 - Charged by Kinesis Processing Units (KPU's) consumed per hour
 - 1 KPU = 1 vCPU + 4GB
- Serverless; scales automatically
- Use IAM permissions to access streaming source and destination(s)
- Schema discovery

we can also use anomaly detection

Kinesis Analytics

- Pay only for resources consumed (but it's not cheap)
 - Charged by Kinesis Processing Units (KPU's) consumed per hour
 - 1 KPU = 1 vCPU + 4GB
- Serverless; scales automatically
- Use IAM permissions to access streaming source and destination(s)
- Schema discovery