

Machine learning aws


Amazon Machine Learning (ML)

- Provides visualization tools & wizards to make creating a model easy
- You point it to training data in S3, Redshift, or RDS
- It builds a model that can make predictions using batches or a low-latency API
- Can do train/test and evaluate your model
- Fully managed
- Honestly it's a bit outdated now



"Ideal Usage Patterns"

- Flag suspicious transactions (fraud detection)
- Forecasting product demand
- Personalization – predict items a user will be interested in
- Predict user activity (we'll do this)
- Classify social media (does this Tweet require my attention?)



Amazon ML: Promises & Limitations

- No downtime
- Up to 100GB training data (more via support ticket)
- Up to 5 simultaneous jobs (more via support ticket)



Amazon ML: Anti-Patterns

- Terabyte-scale data
- Unsupported learning tasks
 - Sequence prediction
 - Unsupervised clustering
 - Deep learning
- EMR / Spark is an (unmanaged) alternative.